

Enhancing Sarcasm Detection using GAN-BERT with Multi-Task Learning

Afif Hossain Irfan¹, Shrabani Das², Jannatul Ferdous³, and Md. Tabil Ahammed⁴

¹Department of CSE, International Islamic University Chittagong (IIUC), Chittagong, Bangladesh

²Department of CSE, Bangladesh University of Business and Technology (BUBT), Dhaka, Bangladesh

³Department of CSE, Chittagong University of Engineering and Technology (CUET), Chattogram, Bangladesh

⁴Department of EEE, Bangladesh University of Engineering and Technology (BUET), Dhaka, Bangladesh

Email: ¹afifhossain.cse.1012@gmail.com, ²shrabanidas99.cse@gmail.com,

³jannatulferdauschow@gmail.com, ⁴ahtabil53@gmail.com

Abstract—Sarcasm detection remains a complex task in natural language processing due to its nuanced and context-dependent nature. This paper introduces a novel sarcasm detection framework combining GAN-BERT and Multitask Learning (MTL), where sentiment and emotion serve as auxiliary supervision tasks to enhance feature representation. The proposed model uses sentiment and emotion as auxiliary tasks to improve its understanding of sarcastic cues. The Sarcasm Headlines Dataset, comprising 28,619 labeled headlines, is augmented with sentiment and emotion annotations utilizing pretrained NLP models. GAN-BERT uses 70% of the 80% training split as labeled data and the remaining 30% as unlabeled, allowing semi-supervised learning. Experimental results demonstrate that the proposed model outperforms BERT, LSTM, and CNN, achieving an accuracy of 93.88% and an F1 score of 94.15%. These findings demonstrate the effectiveness of integrating adversarial training and multitask objectives in sarcasm detection.

Index Terms—Sarcasm Detection, GAN-BERT, Multi-Task Learning, Sentiment Analysis, Emotion Classification, Semi-Supervised Learning, Transformers

I. INTRODUCTION

Sarcasm is one of the most complex linguistic devices; what the speaker says is usually the opposite of the literal meaning of the words used [1]. This implicit contrast renders sarcasm detection difficult in Natural Language Processing (NLP). Compared to fundamental sentiment analysis, sarcasm typically harbors fine-grained contextual hints, incongruity, and affective undertones, which make it very difficult for models without deep semantic understanding. Recent advances in transformer-based models, such as BERT, have achieved staggering performance gains on a broad spectrum of text classification problems through large-scale pretraining and fine-tuning [2], [3]. However, these models tend to require significant quantities of annotated data to achieve good performance. This constitutes a limitation in domains such as sarcasm detection, where annotated sets are small and qualitative. Recent advances in deep learning have shown promise in solving these challenges. Traditional models, such as Long Short-Term Memory (LSTM) and Convolutional Neural Networks (CNN), have successfully acquired pattern-based sequential features in text [4], [5].

Recently, transformer-based models such as BERT have set new benchmarks in NLP by learning rich contextual relationships [6]–[9]. Several studies have highlighted different approaches and challenges in text-based emotion and sarcasm detection, emphasizing the importance of context and the role of big data and deep learning frameworks [10]–[13]. The application of machine learning and deep learning to pattern recognition in textual data has also been extensively investigated in various contexts, including sentiment analysis and context-aware detection systems [14]–[16].

Building upon these foundational developments, various models have been proposed in the literature to enhance sarcasm detection using structured datasets and deep learning techniques. Several sarcasm detection models have been proposed using structured datasets [17]–[31]. A popular benchmark dataset includes 28,619 news headlines, 14,000 sarcastic from The Onion, and 14,619 non-sarcastic from HuffPost [17]. This dataset achieved 77% word embedding coverage, enabling hybrid neural networks, such as CNN and BiLSTM with attention mechanisms, to reach an accuracy of 89.7%. GMP-LSTM achieved 92.54% accuracy and 91.27% F1 score [18]. Traditional approaches like SVM reached 87.5% accuracy, while LSTM showed 90.6% and CNN 93.2% [19], [20]. BERT-based models showed strong performance with Hindi-English (Hinglish) data sets (98.9% precision, 0.98 F1 score) and Urdu tweets (79.51% accuracy using BERT-LSTM) [21]. BERT, using pre-trained embeddings, outperformed models with GloVe and Word2Vec by 7–10% [22]. Context-based models, such as RoBERTa and DistilBERT, reduced training time and utilized metadata, including article section or speaker identity, to enhance performance [23]. A CNN-BiGRU ensemble reached 96.2% precision using BERT and XLNet embeddings [24]. GAN-based models such as GAN-BERT outperformed standard BERT in semisupervised settings, while GAN-CITE used contrastive learning and transformers to perform well in low-resource sarcastic data [25], [26]. Other variations of GAN, such as C-GAN and LC-GAN, used sentiment embeddings, word-level loss, and attention to improve sarcasm detection [27], [28]. RoBERTa, BiLSTM, and

multihead attention combinations achieved 93.36% precision and 93.34% F1 score [29]. Recent trends focus on deep and hybrid models, including transformer-based approaches, attention mechanisms, sentiment features, and GANs, to achieve high accuracy in diverse languages and domains [30], [31]. Despite the development of various methodologies for sarcasm detection, a significant gap remains in comparative analysis between deep learning models, such as CNN and LSTM, and transformer-based models, such as BERT. This gap becomes even more prominent when considering the potential of semi-supervised techniques such as GAN-BERT. Existing studies often focus on individual model performance but lack a systematic comparison, particularly in the context of short-form texts and low-resource scenarios where annotated data are limited.

To address this gap, this study aims to provide a comprehensive comparative review of four deep learning models for sarcasm detection: GAN-BERT, a semi-supervised transformer-based model; BERT, a fully supervised transformer model; LSTM, a sequential deep learning model; and CNN, a convolutional model optimized for local feature detection. All models will be evaluated using the same benchmark dataset of 28,619 labeled news headlines and will undergo consistent preprocessing and training procedures. Their performance will be assessed using standardized evaluation metrics, including accuracy, F1 score, confusion matrix, ROC curves, and AUC values. This comparative analysis aims to identify the most effective model for sarcasm detection, particularly in low-resource conditions, thus informing the selection of practical solutions for real-world applications.

II. PROPOSED METHODOLOGY

The architecture, training procedure, and comparative baseline models utilized in this investigation are described in this section. This work compares the performance of the proposed new sarcasm detection model, which is based on GAN-BERT with Multi-Task Learning (MTL), with three popular baseline models: CNN, LSTM, and BERT in Figure 1.

A. GAN-BERT with Multi-Task Learning (Proposed Model)

By combining the GAN-BERT architecture with Multi-Task Learning (MTL), the proposed model in this paper enhances the detection of sarcasm in brief texts. The central component of the model is a shared BERT encoder that uses input headlines to produce deep contextual embeddings. These embeddings are then fed into three output heads, each responsible for a specific task. The architecture of our proposed model is illustrated in Figure 2. The model integrates a shared BERT encoder with task-specific heads for sarcasm, sentiment, and emotion classification, along with a discriminator for semi-supervised learning. A sarcasm classifier, a binary classification head, performs the main task of identifying whether a headline is sarcastic (1) or non-sarcastic (0). To enhance




Fig. 1: Proposed Sarcasm Detection Workflow using GAN-BERT with Multi-Task Transformer Learning Models

feature learning, two auxiliary tasks are added during training: an emotion classifier that recognizes the underlying emotional tone from several classes, including joy, anger, and disgust, and a sentiment classifier that performs binary classification of the headline's sentiment as positive or negative. Additionally, the model utilizes a GAN-BERT architecture with an adversarial discriminator that learns to distinguish between labeled and unlabeled data. This enables semi-supervised learning, which enhances the model's capacity for generalization, particularly in situations where labeled data is scarce.

1) *Loss Function*: The total training loss for our *GAN-BERT with Multi-Task Learning* model is defined as:

$$L_{\text{total}} = L_{\text{sarcasm}} + \lambda_1 \cdot L_{\text{sentiment}} + \lambda_2 \cdot L_{\text{emotion}} + L_{\text{adv}} \quad (1)$$

To maximize performance, the model's overall training objective integrates several loss components. In particular, $L_{\text{sentiment}}$ and L_{emotion} represent auxiliary losses for sentiment and emotion classification, respectively, while L_{sarcasm} denotes




Fig. 2: GAN-BERT with MTL

the cross-entropy loss for the primary task of sarcasm classification. The hyperparameters λ_1 and λ_2 control the contribution of these auxiliary tasks to the overall loss, thus regulating their influence. Furthermore, L_{adv} represents the adversarial loss obtained from a discriminator, which is trained in a semi-supervised environment to differentiate between labeled and unlabeled samples. Crucially, only the sarcasm classification head is activated during inference. Deactivating the sentiment and emotion classification heads enables effective prediction while retaining the advantages of richer, multi-task supervision during training.

B. BERT (Baseline Transformer Model)

Bidirectional Encoder Representations from Transformers, or BERT, is a potent pre-trained transformer model that excels at recognizing bidirectional context in text. This work uses the `bert base-uncased` model, which has been specifically optimized for binary sarcasm classification, as the baseline of this work. This architecture utilizes the embedding corresponding to the [CLS] token for classification after the headlines are first tokenized using the BERT tokenizer. The final output is obtained by applying a sigmoid activation function after a fully connected dense layer was applied to this embedding. Eighty percent of the dataset is used to train the model using binary cross-entropy loss, with the remaining twenty percent set aside for evaluation. Because of its deep bidirectional context learning capabilities, BERT demonstrated strong performance in sarcasm classification, achieving an accuracy of 92.04%.

C. LSTM (Baseline Sequential Model)

Because they can learn long-range dependencies, Long Short-Term Memory (LSTM), a type of Recurrent Neural Network (RNN), is ideal for modeling sequential data, such as sentences. The architecture first tokenizes and transforms input

headlines into padded sequences. The words are represented by a pre-trained GloVe embedding layer using 100-dimensional vectors. A unidirectional LSTM layer with 128 hidden units, a dropout layer to avoid overfitting, and a final dense layer with a sigmoid activation function for binary classification come next. Binary cross-entropy loss is used to train the model on 80% of the dataset. The LSTM model achieved an accuracy of 87.16% in terms of performance. Compared to transformer-based models, it performs poorly in contextual understanding despite being good at capturing sequential patterns.

D. CNN (Baseline N-Gram Feature Model)

Although primarily developed for image data, Convolutional Neural Networks (CNNs) have proven highly effective in text classification tasks due to their ability to capture local n-gram patterns. The architecture represents input text using a pre-trained GloVe embedding layer. To extract local features from embedded sequences, several 1D convolutional filters with kernel sizes of 3, 4, and 5 are used. To preserve the most noticeable features, the outputs of these convolution layers are subsequently combined using global max pooling. For binary classification, these pooled representations are concatenated and fed into a dense layer that has a sigmoid activation function. Binary cross-entropy loss is used to train the model on 80% of the dataset. The CNN model performed well, achieving an accuracy of 85.18%. In tasks like sarcasm detection, which frequently require nuanced and context-rich language, its limited capacity to capture broader context diminishes its efficacy, despite its speed and computational simplicity.

E. Summary

With the highest accuracy and F1 score of any model, GAN-BERT, using the MTL framework, performed the best overall. This demonstrates that a more contextually aware and broadly applicable sarcasm detection model

is produced by combining semi-supervised learning with multitask objectives. The general workflow of our system, from the augmentation of the data set to the final prediction, is presented in Figure 1. It summarizes the training pipeline and highlights the semi-supervised setup used in GAN-BERT with Multi-Task Learning.

III. DATASET AND PREPROCESSING

The *Sarcasm Headlines Dataset*, obtained from Kaggle [32], was used in this investigation. It is ideal for sarcasm detection tasks because it includes 28,619 samples in JSON format. Each dataset entry has three fields: "article_link" that contains the URL of the whole article (not used in this study); "is_sarcastic" that includes a binary target label indicating sarcasm (1 for sarcastic, 0 for non-sarcastic); and "headline" that contains the brief news headline used as input text. Using headlines from two different sources, *The Onion*, well-known for its satirical content, and *HuffPost*, which is not, this dataset provides a formal and varied text corpus that effectively conveys the complex language frequently associated with sarcasm.

A. Preprocessing Steps and Data Splitting Strategy

All headline text was cleaned up by deleting punctuation, special characters, URLs, and non-ASCII characters as part of the preprocessing step. By doing this, the input text was ensured to be consistent and free from noise that could impair the model's performance. Depending on the model type, two data splitting strategies were used to enable impartial and consistent evaluation across all models. A typical 80/20 split was used for CNN, LSTM, and BERT models, with 20% (5,724 samples) reserved for testing and 80% of the dataset (22,895 samples) used for training. The GAN-BERT model with Multi-Task Learning (MTL), on the other hand, began with the same 80/20 train-test split, but separated the 80% training portion into 30% unlabeled data (approximately 6,868 samples) and 70% labeled data (approximately 16,026 samples). This setup is designed to facilitate semi-supervised adversarial training, where the generator (BERT encoder) enhances its representation learning through task-specific supervision and adversarial feedback. In contrast, the discriminator learns to distinguish between labeled and unlabeled inputs.

B. Data Augmentation for Multi-Task Learning

To facilitate efficient multitask learning, additional task labels for sentiment and emotion were automatically generated, as the original dataset only contained binary sarcasm labels. A pre-trained model called *cardiffnlp/twitter-roberta-base-sentiment* was used for Sentiment Labels. This model distinguishes three types of text: *positive*, *neutral*, and *negative*. The neutral class was eliminated, leaving only the positive and negative classes, thereby reducing the task to a binary auxiliary objective. One of two reliable sources was used to derive annotations for Emotion Labels. The first is Google's GoEmotions model, which was trained using a

wide range of 28 emotion classes, including sadness, anger, and joy. The PySentimiento library is a Python toolkit for affective natural language processing, providing fine-grained emotion classification. The output of both approaches was mapped into more general, high-level emotion categories, including *Surprise*, *Joy*, *Anger*, *Disgust*, and *Sadness*. The dataset was enhanced by this augmentation process, providing the model with more varied and complex training data.

Five example entries from the Sarcasm Headlines Dataset after data augmentation with sentiment and emotion labels using pre-trained models are shown in Table I.

TABLE I: SNIPPET OF DATASET AFTER AUGMENTED IN SARCASM HEADLINES DATASET

Headline	Sarcasm	Sentiment	Emotion
"Another genius plan by Congress"	1	Negative	Disgust
"Vaccines cause autism, says no scientist ever"	1	Positive	Anger
"Stock prices rise again amid investor optimism"	0	Positive	Joy
"Can't wait for Monday morning traffic"	1	Negative	Frustration
"Local man wins lottery after losing job last week"	0	Positive	Surprise

Notes [Table I]:

Sarcasm: Original binary label from the dataset (1 = sarcastic, 0 = not sarcastic)

Sentiment: Generated using *cardiffnlp/twitter-roberta-base-sentiment* (positive/negative)

Emotion: Generated using GoEmotions or *pysentimiento* (mapped to major classes such as joy, anger, etc.)

IV. RESULTS AND DISCUSSION

A. Classification Performance

In this paper, the performance of the suggested GAN-BERT model is compared with that of the Multi-Task Learning (MTL) model and the conventional baselines, including CNN, LSTM and BERT, to assess its efficacy. An 80/20 split of the Sarcasm Headlines Dataset was used to train and test the models. To enable semi-supervised learning, 70% of the training data for GAN-BERT was labeled, and 30% was unlabeled. The classification report is summarized in Table II.

As shown in Table II, GAN-BERT with MTL achieves the highest accuracy and F1 score, outperforming all baseline models including standard BERT, LSTM and CNN. The proposed model outperformed all baselines by a significant margin, achieving the highest accuracy and F1 score, particularly for sarcastic texts. This demonstrates how the encoder can become more generalizable by including auxiliary tasks (sentiment and emotion).

B. Accuracy and Loss Curves

Figures 3 through 6 present the training and validation accuracy and loss plots for each model evaluated in this study:

TABLE II: DETAILED CLASSIFICATION REPORT ACROSS MODELS

Model	Class	Precision	Recall	F1-Score	Support
GAN-BERT with MTL	Non-sarcastic (0)	0.9412	0.9307	0.9359	14,149
	Sarcastic (1)	0.9384	0.9448	0.9415	14,470
	Accuracy			0.9388	28,619
	Macro Avg	0.9398	0.9377	0.9387	
	Weighted Avg	0.9399	0.9388	0.9387	
BERT	Non-sarcastic (0)	0.9123	0.9245	0.9183	14,149
	Sarcastic (1)	0.9310	0.9136	0.9222	14,470
	Accuracy			0.9204	28,619
	Macro Avg	0.9216	0.9190	0.9202	
	Weighted Avg	0.9207	0.9204	0.9203	
LSTM	Non-sarcastic (0)	0.8642	0.8727	0.8684	14,149
	Sarcastic (1)	0.8791	0.8702	0.8746	14,470
	Accuracy			0.8716	28,619
	Macro Avg	0.8717	0.8715	0.8715	
	Weighted Avg	0.8716	0.8716	0.8716	
CNN	Non-sarcastic (0)	0.8421	0.8539	0.8479	14,149
	Sarcastic (1)	0.8613	0.8495	0.8553	14,470
	Accuracy			0.8518	28,619
	Macro Avg	0.8517	0.8517	0.8516	
	Weighted Avg	0.8518	0.8518	0.8518	

the proposed GAN-BERT with MTL, BERT, LSTM, and CNN. As shown in Figure 3, the GAN-BERT with MTL model demonstrates smoother convergence and improved generalization compared to other models. This performance trend is consistent with the observations in Figure 4 through Figure 6. With little overfitting, the plots exhibit steady convergence over training epochs. Both the auxiliary task supervision and the adversarial loss component contribute to the successful learning dynamics of the model.




Fig. 3: Accuracy and Loss Curve of GAN-BERT with MTL

C. Comparative Analysis with Existing Work

Table III summarizes the performance of the proposed model against recent studies on sarcasm detection. The proposed GAN-BERT + MTL outperforms previous models in terms of both accuracy and F1 score, showcasing the benefit of multitask and semi-supervised learning.

D. Discussion

The results of the study validate the main hypothesis, which states that combining multitask learning with semisupervised adversarial training significantly enhances sarcasm detection




Fig. 4: Accuracy and Loss Curve of BERT




Fig. 5: Accuracy and Loss Curve of LSTM




Fig. 6: Accuracy and Loss Curve of CNN

TABLE III: COMPARISON OF PROPOSED WORK WITH PREVIOUS STUDIES

Study	Model / Method	Dataset	Accuracy
Proposed Work (2025)	GAN-BERT + MTL	News Headlines (Kaggle)	0.9388
Khan et al. (2025)	RoBERTa + BiLSTM + MHA	News Headlines	0.9336
Ali et al. (2023)	GMP-LSTM (Global Max Pooling + LSTM)	Onion & HuffPost + Corpus V2	0.9254
Sharma et al. (2023)	Word2Vec/GloVe/BERT + Fuzzy Logic Ensemble	News Headlines, SARC, Twitter	0.9081*
Croce et al. (2020)	GAN-BERT	Multi-benchmark datasets	~0.90–0.94

* Accuracy reported on *News Headlines* subset only (from Sharma et al., 2023).

in short texts. By including sentiment and emotion as auxiliary tasks, the model can capture more complex linguistic patterns that transcend surface-level textual features. One of the key findings is that the F1 score for sarcasm detection improved from 0.9222 with the baseline BERT model to 0.9415 with the GAN-BERT + MTL approach. Additionally, the macro F1 score and weighted average precision demonstrate strong and balanced performance in both the sarcastic and non-sarcastic classes. Importantly, the semi-supervised learning framework allowed the model to use unlabeled data, which is especially helpful in low-resource NLP scenarios where there aren't enough labeled examples.

V. CONCLUSION

By combining GAN-BERT with Multi-Task Learning (MTL) and utilizing sentiment and emotion classification as auxiliary tasks, this study presents a novel approach to sarcasm detection. By using a semi-supervised adversarial framework that utilizes both labeled and unlabeled data, the model enhances its ability to recognize subtle patterns in brief, sarcastic texts. The proposed model outperforms conventional architectures, such as BERT, LSTM, and CNN, achieving an accuracy of 93.88% and an F1 score of 94.15% on the Sarcasm Headlines dataset. The best performance was specifically obtained by GAN-BERT with MTL, demonstrating the value of adversarial training and multitask objectives in enhancing sarcasm classification. This study demonstrates that incorporating auxiliary signals, such as sentiment polarity and emotional tone, improves feature learning and generalization in low-resource settings. Future directions include testing this framework on more informal domains, such as social media data, investigating humor and irony as additional auxiliary tasks, and applying it to multilingual sarcasm detection.

REFERENCES

- [1] B. Yao, Y. Zhang, Q. Li, and J. Qin, "Is sarcasm detection a step-by-step reasoning process in large language models?," *Proc. Conf. AAAI Artif. Intell.*, vol. 39, no. 24, pp. 25651–25659, 2025.
- [2] M. T. Ahammed et al., "Survey on wireless fidelity (Wi-Fi) based indoor location technology for mobile devices," in *Proc. 3rd Int. Conf. Intell. Comput. Instrum. Control Technol. (ICICICT)*, 2022, pp. 1465–1470.
- [3] P. Chandra et al., "Contextual emotion detection in text using deep learning and big data," in *Proc. 2nd Int. Conf. Comput. Sci., Eng. Appl. (ICCSEA)*, 2022, pp. 1–5.
- [4] S. Deena et al., "Real-time based Violence Detection from CCTV Camera using Machine Learning Method," in *Proc. Int. Conf. Industry 4.0 Technol. (I4Tech)*, vol. 2, no. 1, 2022, pp. 1–6.
- [5] M. T. Ahammed, S. Ghosh, and M. A. R. Ashik, "Human and object detection using machine learning algorithm," in *Proc. Trends Electr., Electron., Comput. Eng. Conf. (TEECCON)*, 2022, pp. 39–44.
- [6] M. T. Ahammed, M. A. Rahman, and S. Ghosh, "Surface plasmon resonance sensor based on Ag layer coated PCF for dry sandy soil detection with deep learning algorithm," in *Proc. 5th Int. Conf. Sustainable Technol. Industry 5.0 (STI)*, 2023, pp. 1–6.
- [7] M. Ibrahim et al., "Design a novel top as based suspended core single mode photonic crystal fiber optics communications," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Informatics (ACCAI)*, vol. 64, no. 12, 2022, pp. 1–5.
- [8] M. Ibrahim et al., "Designing a top as based suspended square core low loss single-mode photonic crystal fiber optics communications," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Informatics (ACCAI)*, vol. 8, 2022, pp. 1–6.
- [9] M. T. Ahammed and M. F. Hossain, "Inverse design of a PCF-based SPR sensor using machine learning," in *Proc. IEEE Int. Conf. Comput., Appl. Syst. (COMPAS)*, 2024, pp. 1–6.
- [10] T. Nisat et al., "Big data analysis for E-trading flower shop management system," in *Proc. Int. Conf. Commun., Comput. Internet Things (IC3IoT)*, vol. 10, no. 2, 2022, pp. 1–6.
- [11] S. Ghosh et al., "Researching the influence of peer-reviews in action movies based on public opinion," *J. Media, Culture and Communication*, vol. 2, no. 23, pp. 8–13, 2022.
- [12] M. T. Ahammed et al., "Sentiment analysis using a machine learning approach in python," in *Proc. Int. Conf. Commun., Comput. Internet Things (IC3IoT)*, 2022, pp. 1–6.
- [13] S. Prathibha et al., "Detection methods for software defined networking intrusions (SDN)," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Informatics (ACCAI)*, 2022, pp. 1–6.
- [14] M. M. Mahmud, S. Das, F. A. Adhara, and M. T. Ahammed, "PCF-based SPR sensor for alcohol detection," in *Proc. Int. Conf. Recent Prog. Sci., Eng. Technol. (ICRPSET)*, 2024, pp. 1–5.
- [15] C. G. Kousalya et al., "Design of IoT based indoor air purifier," in *Proc. 3rd Int. Conf. Intell. Comput. Instrum. Control Technol. (ICICICT)*, 2022, pp. 1534–1539.
- [16] A. Kabir et al., "Light fidelity (Li-fi) based indoor communication system," in *Proc. Int. Conf. Adv. Comput., Commun. Appl. Informatics (ACCAI)*, vol. 53, no. 10, 2022, pp. 1–5.
- [17] R. Misra and P. Arora, "Sarcasm detection using news headlines dataset," *AI Open*, vol. 4, pp. 13–18, 2023.
- [18] A. Khan, D. Majumdar, and B. Mondal, "A hybrid transformer based model for sarcasm detection from news headlines," *J. Intell. Inf. Syst.*, 2025.
- [19] R. Ali et al., "Deep learning for sarcasm identification in news headlines," *Appl. Sci. (Basel)*, vol. 13, no. 9, p. 5586, 2023.
- [20] M. E. Hassan et al., "Detection of sarcasm in Urdu tweets using deep learning and transformer based hybrid approaches," *IEEE Access*, vol. 12, pp. 61542–61555, 2024.
- [21] P. Shrikhande, V. Setty, and D. A. Sahani, "Sarcasm detection in newspaper headlines," in *Proc. IEEE Int. Conf. Ind. Inf. Syst. (ICIIS)*, 2020, pp. 483–487.
- [22] D. K. Nayak and B. K. Bolla, "Efficient deep learning methods for sarcasm detection of news headlines," in *Machine Learning and Autonomous Systems*, Singapore: Springer, 2022, pp. 371–382.
- [23] K. Chatrinan, T. Noraset, and S. Tuarob, "GAN-CITE: leveraging semi-supervised generative adversarial networks for citation function classification with limited data," *Scientometrics*, vol. 130, no. 2, pp. 679–703, 2025.
- [24] D. K. Sharma et al., "Sarcasm detection over social media platforms using hybrid ensemble model with fuzzy logic," *Electronics (Basel)*, vol. 12, no. 4, p. 937, 2023.
- [25] P. Goel et al., "Sarcasm detection using deep learning and ensemble learning," *Multimed. Tools Appl.*, vol. 81, no. 30, pp. 43229–43252, 2022.
- [26] N. A. Helal et al., "A contextual-based approach for sarcasm detection," *Sci. Rep.*, vol. 14, no. 1, p. 15415, 2024.
- [27] O. Prakash, "Sarcasm detection for sentiment analysis using deep learning enhancement," *Procedia Comput. Sci.*, vol. 258, pp. 203–212, 2025.
- [28] C. I. Eke, A. A. Norman, and L. Shuib, "Context-based feature technique for sarcasm identification in benchmark datasets using deep learning and BERT model," *IEEE Access*, vol. 9, pp. 48501–48518, 2021.
- [29] R. Pandey and J. P. Singh, "BERT-LSTM model for sarcasm detection in code-mixed social media post," *J. Intell. Inf. Syst.*, 2022.
- [30] G. Yu et al., "Adversarial active learning for the identification of medical concepts and annotation inconsistency," *J. Biomed. Inform.*, vol. 108, p. 103481, 2020.
- [31] D. Croce, G. Castellucci, and R. Basili, "GAN-BERT: Generative adversarial learning for robust text classification with a bunch of labeled examples," in *Proc. 58th Annu. Meet. Assoc. Comput. Linguist.*, 2020, pp. 2114–2119.
- [32] R. Misra, "News Headlines Dataset for Sarcasm Detection," *Kaggle*, 2022. [Online]. Available: <https://www.kaggle.com/datasets/rmisra/news-headlines-dataset-for-sarcasm-detection> [Accessed: Apr. 10, 2025].